

操作系统研讨课

Course: B0911011Y

蒋德钧

Fall Term 2023-2024

email: jiangdejun@ict.ac.cn

office phone: 62601007



Lecture 2 A Simple Kernel (Part I)

2023.09.25




Schedule

- Project 1 due
- Project 2 assignment



Project 2 – A Simple Kernel

- Booting procedure



Project 2 – A Simple Kernel

- Requirements (part I)
 - Write a simple kernel (non-preemptive)
 - Start a set of processes
 - Perform context switches between processes
 - Provide non-preemptive kernel support with context switch
 - Support basic mutex to allow BLOCK state of processes



Project 2 – A Simple Kernel

- A set of user processes
 - Program codes under the *test/test_project2* directory in start code
 - STRONGLY suggest to first read the codes of different tasks to understand what they do
 - You need to initialize these tasks in `init_task_info(main.c)` as you did in Project 1



Project 2 – A Simple Kernel

- Process Control Block (PCB)
 - A in-memory data structure in OS kernel containing the information to manage a process
 - Please refer to the given definition `pcb_t` (*include/os/sched.h*)



Project 2 – A Simple Kernel

- Process Control Block (PCB)
 - Process ID
 - Process status
 - Kernel stack pointer
 - User stack pointer
 - You may need to add new fields to PCB in subsequent projects

```
/* Process Control Block */
typedef struct pcb
{
    /* register context */
    // NOTE: this order must be preserved, which is defined in regs.h!!
    reg_t kernel_sp;
    reg_t user_sp;

    /* previous, next pointer */
    list_node_t list;

    /* process id */
    pid_t pid;

    /* BLOCK | READY | RUNNING */
    task_status_t status;

    /* cursor position */
    int cursor_x;
    int cursor_y;

    /* time(seconds) to wake up sleeping PCB */
    uint64_t wakeup_time;
} pcb_t;
```



Project 2 – A Simple Kernel

- Initialize PCB

- Initialize PCBs for tested processes
 - Assign ID, set status/cursor_x/cursor_y
 - PCBs are organized as a linked list
- Allocate memory for kernel stack and user stack
 - Please use allocKernelPage/allocUserPage functions (*kernel/mm/mm.c*)

地址范围	建议用途
0x50000000-0x50200000	BBL 代码及其运行所需的内存
0x50200000-0x50500000	Kernel 的数据段/代码段等
0x50500000-0x52000000	供内核动态分配使用的内存
0x52000000-0x52500000	用户程序的数据段/代码段等
<u>0x52500000-0x60000000</u>	<u>供用户动态分配使用的内存</u>

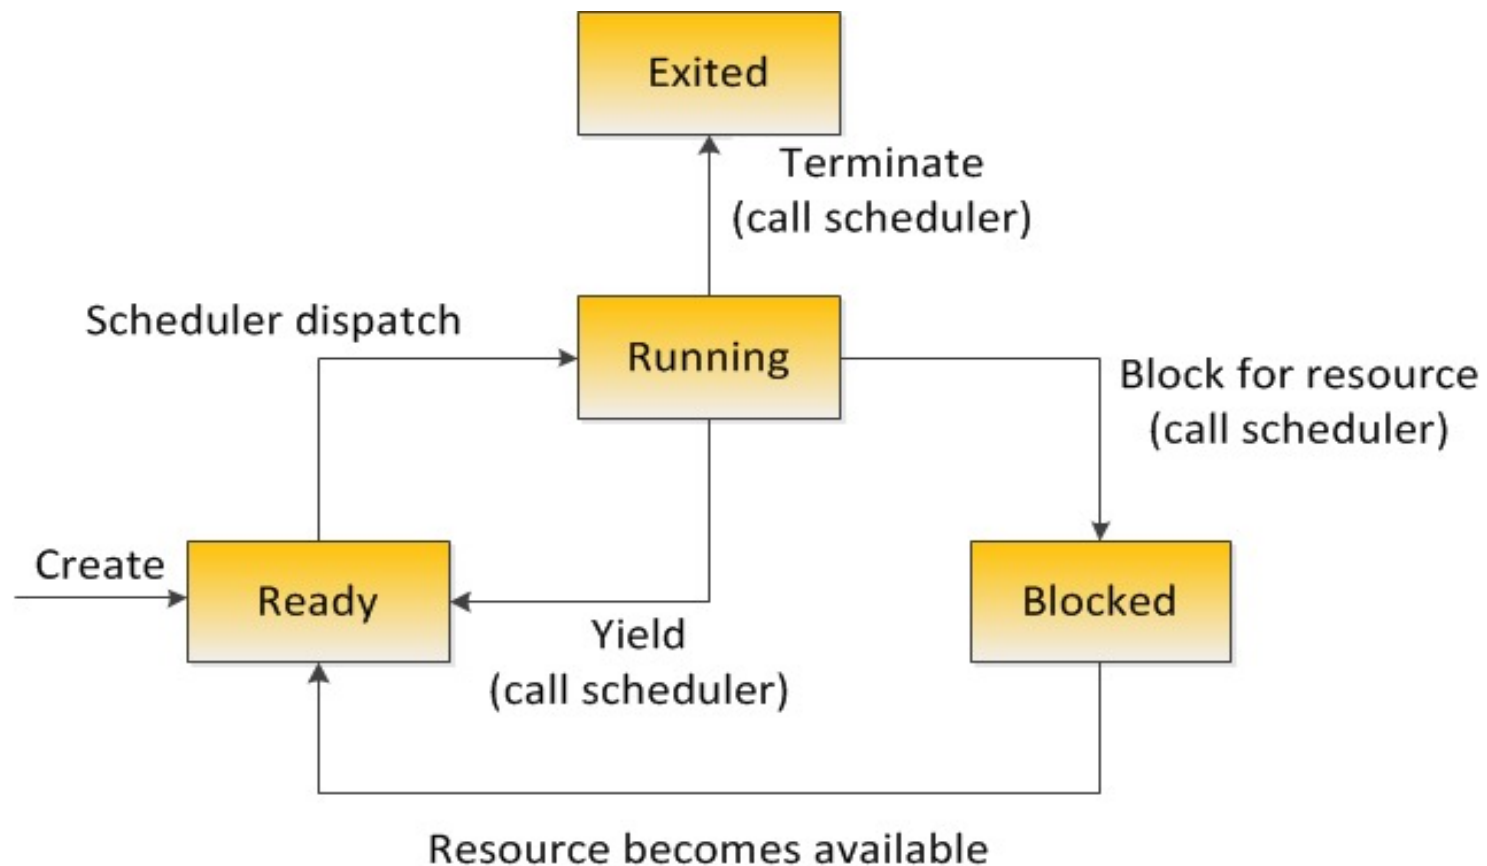
Project 2 – A Simple Kernel

- Initialize kernel stack
 - Kernel stack is used to save and restore the context of a process

```
static void init_pcb_stack(  
    ptr_t kernel_stack, ptr_t user_stack, ptr_t entry_point,  
    pcb_t *pcb)  
{  
    /* TODO: [p2-task3] initialization of registers on kernel stack  
    * HINT: sp, ra, sepc, sstatus  
    * NOTE: To run the task in user mode, you should set corresponding bits  
    *       of sstatus(SPP, SPIE, etc.).  
    */  
    regs_context_t *pt_regs =  
        (regs_context_t *) (kernel_stack - sizeof(regs_context_t));  
  
    /* TODO: [p2-task1] set sp to simulate just returning from switch_to  
    * NOTE: you should prepare a stack, and push some values to  
    * simulate a callee-saved context.  
    */  
    switchto_context_t *pt_switchto =  
        (switchto_context_t *) ((ptr_t) pt_regs - sizeof(switchto_context_t));  
}
```

Project 2 – A Simple Kernel

- Scheduler (non-preemptive kernel)



Project 2 – A Simple Kernel

- Yield
 - A process itself releases the control of CPU by calling kernel scheduler
 - Kernel scheduler places the current running process to the end of the ready queue, and chooses another process to run
 - Please refer to `kernel_yield()` function



Project 2 – A Simple Kernel

- Context switch
 - Kernel scheduler executes context switch to run new process
 - Save context
 - Registers → Memory
 - Restore context
 - Memory → Registers



Project 2 – A Simple Kernel

- switch_to asm function
 - Work like a function call, but after the call the return address is another process
 - You need to save/restore process context in switch_to
 - Where to place the process context?
 - Kernel stack
 - Please refer to switchto_context_t (*include/os/sched.h*) for initializing kernel stack and performing context switch



Project 2 – A Simple Kernel

- Start a process
 - After initialization, the kernel calls `do_scheduler`

```
while (1)
{
    // If you do non-preemptive scheduling, it's used to surrender control
    do_scheduler();

    // If you do preemptive scheduling, they're used to enable CSR_SIE and wfi
    // enable_preempt();
    // asm volatile("wfi");
}
```



Project 2 – A Simple Kernel

- Start a process
 - You need to choose a new process to run in `do_scheduler`
 - Round robin is a preferred scheduling algorithm
 - `switch_to` is called to perform context switch between the current running process and the new one



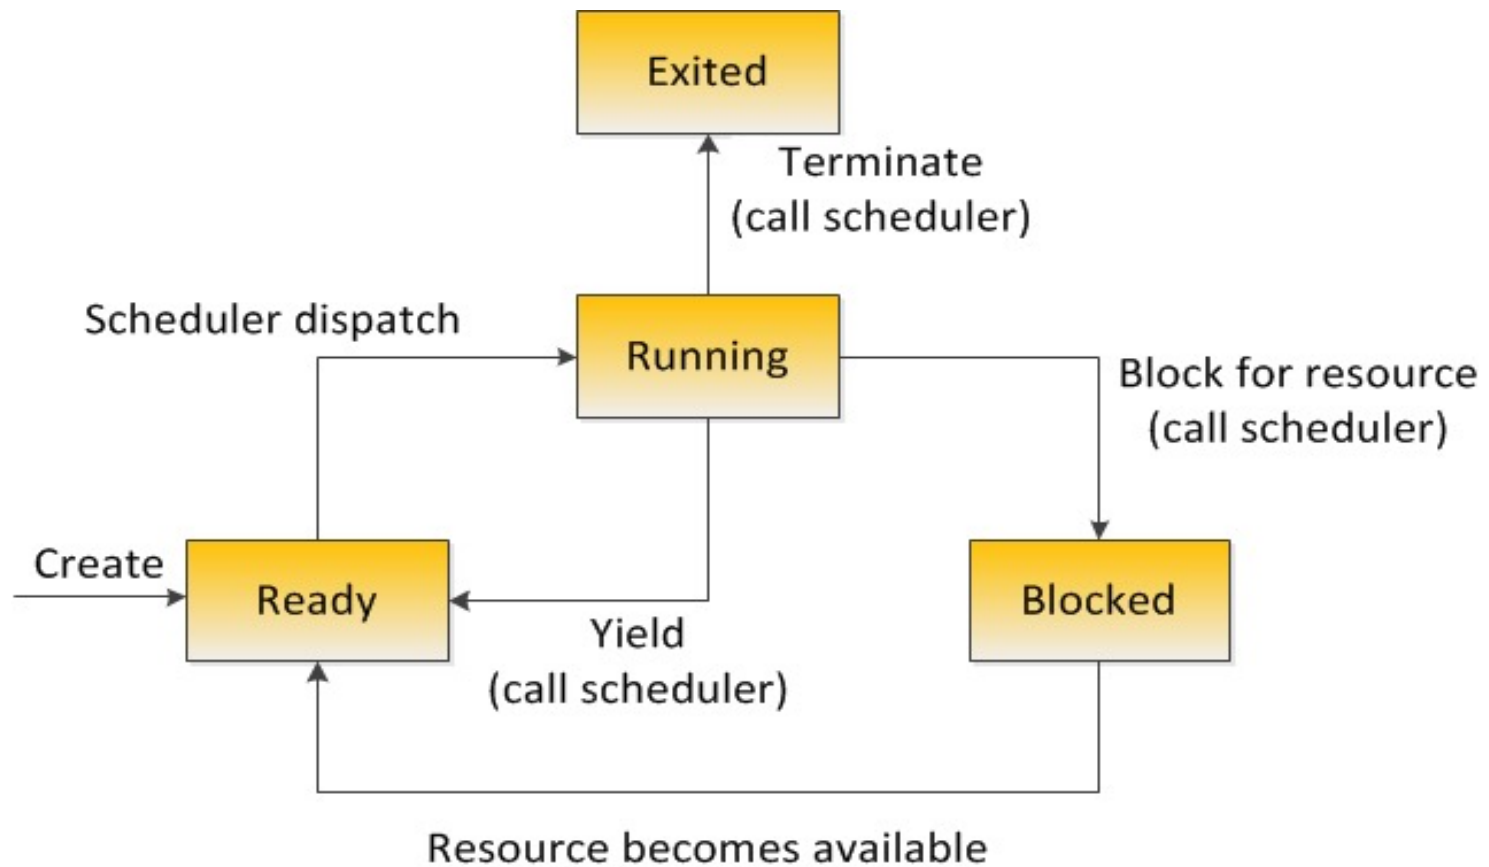
Project 2 – A Simple Kernel

- Start a process
 - Note that, when you start the first process, there is no current running user process, how do you handle this?
 - Please refer to `pid0_pcb` (*include/os/sched.h*)
 - It can be used for store context of kernel process (`main.c`)



Project 2 – A Simple Kernel

- Mutex lock



Project 2 – A Simple Kernel

- Mutex lock
 - What if no process currently holds the lock?
 - Acquire the lock
 - What if the lock is currently held?
 - Wait
 - Implement lock-related functions
 - Manage processes that do not acquire the lock
 - Ready queue vs. wait queue?



Project 2 A Simple Kernel

- Step by step – Task 1
 - Initialize PCBs for tested processes
 - Implement `do_scheduler` and `switch_to`
 - Start tested processes and support context switch among these processes as a non-preemptive kernel



Project 2 A Simple Kernel

- Step by step – Task 2
 - Implement mutex lock to support BLOCK state



Project 2 A Simple Kernel

- Requirement for design review
 - *switch_to*函数的工作机制是怎样的？是如何支持两个进程的上下文切换的？
 - 你在初始化PCB和内核栈时需要做哪些事情？
 - 你在switchto_context_t中会保存哪些内容？
 - 请介绍你设计的内核调度器(do_scheduler)的工作流程



Project 2 A Simple Kernel

- Requirement for design review
 - 当一个进程被阻塞时，内核会进行如何的处理？
 - 当一个进程被阻塞或获得资源时，内核会将进程的PCB放置在哪里？
 - 你设计的互斥锁机制能否支持一个进程请求多把锁？



Project 2 – A Simple Kernel

- Requirement for S/A/C-Core

Core type	Task requirements
S-Core	Tasks 1, 2
A-Core	NA
C-Core	NA



Project 2 – A Simple Kernel

- P2 schedule
 - 9th Oct.
 - P2 part I design review
 - Assign P2 part II
 - 16th Oct.
 - P2 part I due
 - P2 part II design review if it was assigned

